

22. Metrics Monitoring and Alerting System

Real-life examples

- Datadog
- New Relic
- Prometheus
- Graphite

Requirements clarification

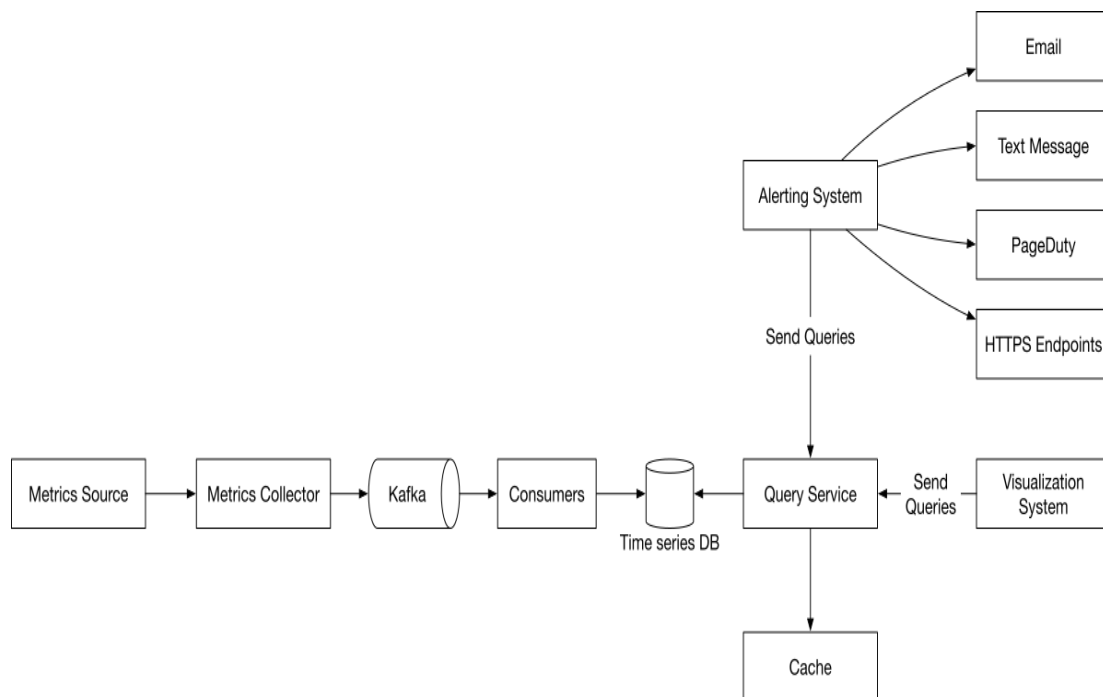
- **Functional requirements**
 - Collect a variety of metrics
 - ✓ CPU usage
 - ✓ Request count
 - ✓ Memory usage
 - ✓ Message count in message queues
 - Send alert notifications to various alerting destinations.
 - ✓ Email
 - ✓ Phone
 - ✓ PagerDuty
 - ✓ Webhooks
 - Support metrics visualization
 - Data retention policy:
 - ✓ Raw form for 7 days.
 - ✓ 1-minute resolution for 30 days.
 - ✓ 1-hour resolution for 1 year.
- **Non-functional requirements**
 - ✓ **Scalability:** The system should be scalable to accommodate growing metrics and alert volume.

- ✓ **Low latency:** The system needs to have low query latency for dashboards and alerts.
- ✓ **Reliability:** The system should be highly reliable to avoid missing critical alerts.
- ✓ **Flexibility:** Technology keeps changing, so the pipeline should be flexible enough to easily integrate new technologies in the future.

System interface definition

Data model definition

High-level design



- **Metrics Source**

- ✓ Provides metrics, like application servers, database, message queues.

- **Metrics Collector**

- ✓ Gathers metrics data and writes data into the time-series database.

- **Kafka**

- ✓ Decouples the data collection and data processing services from each other.

- ✓ Prevents data loss when the database is unavailable, by retaining the data in Kafka.
- ✓ It could be entirely replaced by a competent time-series database (Facebook's Gorilla).
- **Time-series database**
 - ✓ Stores metrics data as time series.
 - ✓ Provides a custom query interface for analyzing and summarizing a large amount of time-series data.
 - ✓ Maintains indexes on labels to facilitate the fast lookup of time-series data by labels.
- **Query Service**
 - ✓ Makes it easy to query and retrieve data from the time-series database.
 - ✓ This should be a very thin wrapper if we choose a good time-series database. It could also be entirely replaced by the time-series database's own query interface.
- **Alerting system**
 - ✓ Sends alert notifications to various alerting destinations.
- **Cache**
 - ✓ Reduces the load of the time-series database and make query service more performant.

Detailed design

Metrics collection

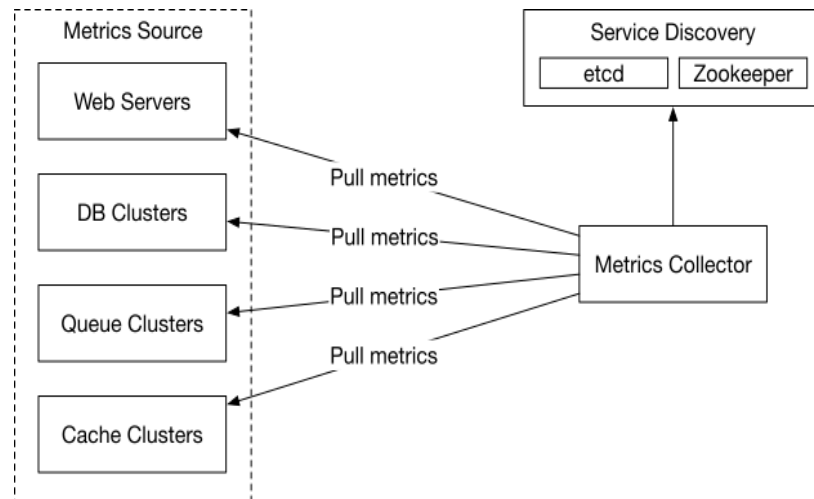
- Ways of collecting metrics

Pull model

There are dedicated metric collectors which pull metrics values from the running applications periodically via a pre-defined HTTP endpoint (for example, /metrics).

Metrics collectors will query Service Discovery to get the complete list of service endpoints to pull data from.

Use a consistent hash ring to decide each metrics collector needs to collect which set of servers.

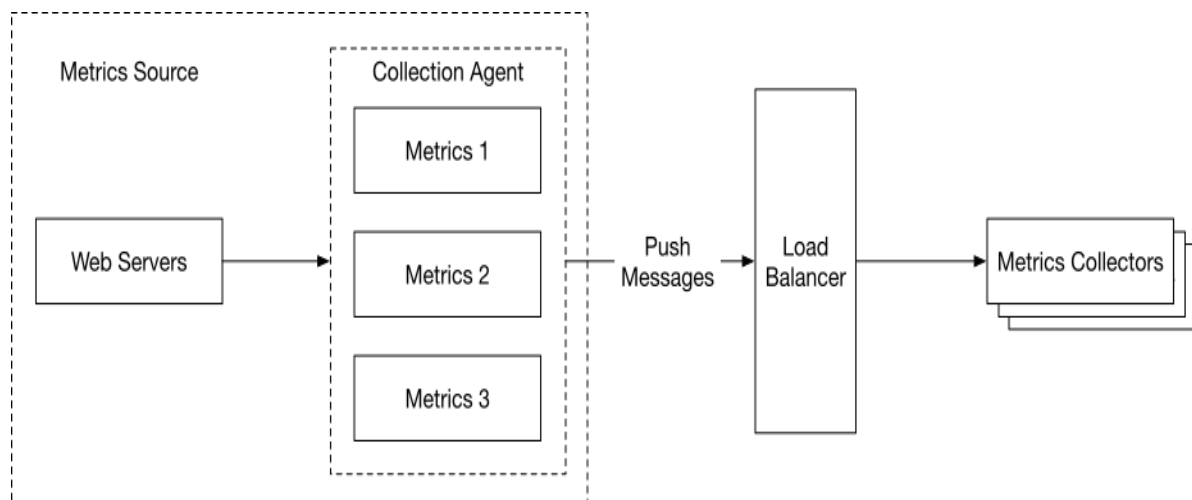


Push model

A collection agent

- Commonly installed on every server being monitored.
- Collects metrics and pushes those metrics periodically to the metrics collector.

The metrics collector should be in an auto-scaling cluster with a load balancer in front of it.

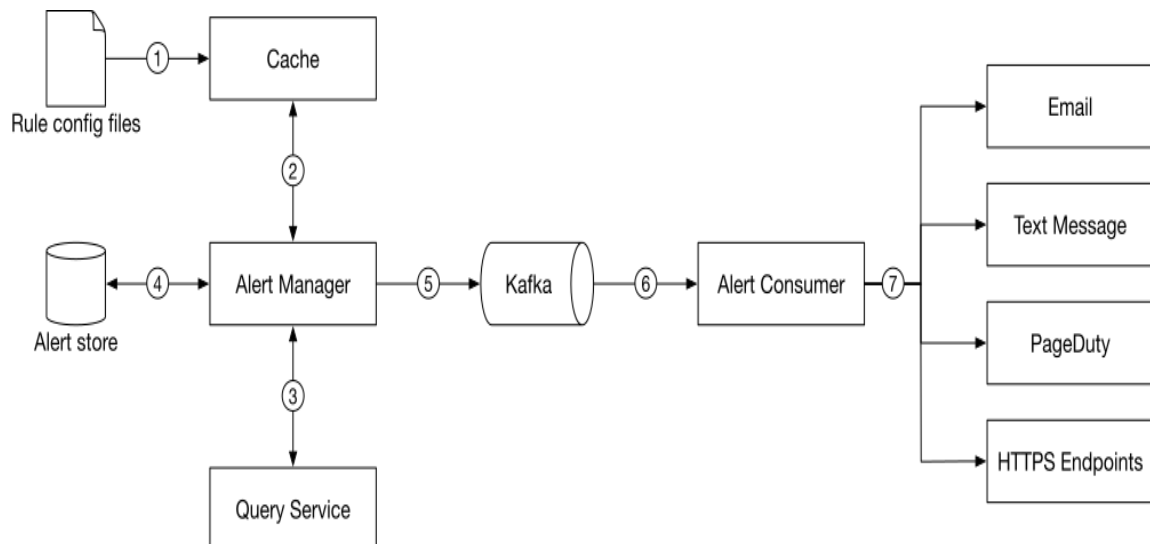


Scale the metrics transmission pipeline

Scale through Kafka

- Configure the number of partitions based on throughput requirements.
- Partition metrics data by metric names, so consumers can aggregate data by metrics names.
- Further partition metrics data with tags/labels.
- Categorize and prioritize metrics so that important metrics can be processed first.

Alerting system



Workflow

- Step 1: Load config files (rules) to cache servers.
- Step 2: The alert manager fetches alert configs from the cache.
- Step 3: Based on config rules, the alert manager calls the query service at a predefined interval.

If the value violates the threshold, an alert event is created.

- Step 4: The alert store is a key-value database and keeps the state (inactive, pending, firing, resolved) of all alerts. It ensures a notification is sent at least once.
- Step 5: Eligible alerts are inserted into Kafka.
- Step 6: Alert consumers pull alert events from Kafka.
- Step 7: Alert consumers process alert events from Kafka and send notifications over to different channels such as email, text message, PagerDuty, or HTTP endpoints.

Notes

- Consider to build or buy alerting system (PagerDuty).

Storage

- Space optimization
 - Data encoding and compression
 - Downsampling (Convert high-resolution data to low-resolution)

Key points

- There are 2 ways to collect metrics: Pull model or Push model.
- The metrics collector cluster should be set up for auto-scaling, to ensure that there are an adequate number of collector instances to handle the demand.
- There are 3 ways to reduce disk usage: Encoding, compression and Downsampling
- Build vs buy options for alerting and visualization systems.

Summary

